



**Fernando Camargos**

Senior Architect

Percona

# Deploying MySQL on Kubernetes with the PerconaOperator



Monday, May 22nd



01:30 PM MST



# Deploying MySQL on Kubernetes with the Percona Operator

Fernando Laudaes Camargos, Senior Architect  
Chetan Shivashankar, Kubernetes Technical Lead

# Table of Contents

---

1. Installation and customization
2. Cluster management
3. Troubleshooting
4. Extras

# Installation and customization

# Installation and customization

---

- Deploying on GKE
- Anti-affinity and tolerations
- Changing MySQL options
- Custom resources options
- Application and system users

# Deploying on GKE

<https://docs.percona.com/percona-operator-for-mysql/pxc/gke.html>

```
git clone -b v1.12.0 https://github.com/percona/percona-xtradb-cluster-operator
```

```
cd percona-xtradb-cluster-operator
```

```
gcloud auth login
```

```
gcloud container clusters create my-cluster-1 --project <project name> --zone us-central1-a --cluster-version 1.23 --machine-type n1-standard-4 --num-nodes=3
```

```
gcloud container clusters get-credentials my-cluster-1 --zone us-central1-a --project <project name>
```

```
kubectl create clusterrolebinding cluster-admin-binding --clusterrole cluster-admin --user $(gcloud config get-value core/account)
```

```
kubectl create namespace pxc
```

```
kubectl config set-context $(kubectl config current-context) --namespace=pxc
```

```
kubectl apply -f deploy/crd.yaml
```

```
kubectl apply -f deploy/rbac.yaml
```

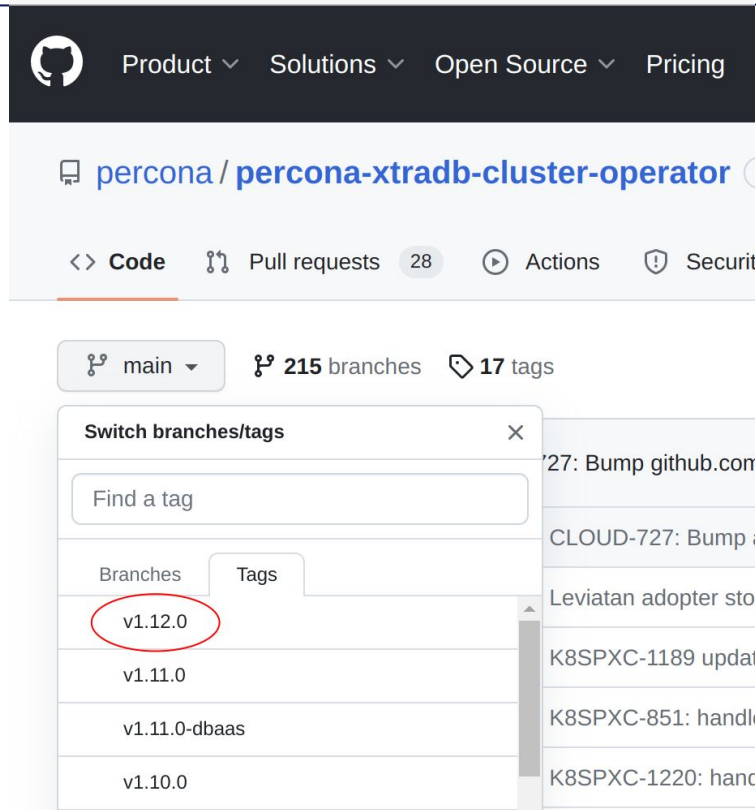
```
kubectl apply -f deploy/operator.yaml
```

```
kubectl create -f deploy/secrets.yaml
```

```
kubectl apply -f deploy/bundle.yaml
```

```
kubectl apply -f deploy/cr.yaml
```

# Deploying on GKE



percona / percona-xtradb-cluster-operator

<> Code Pull requests 28 Actions Security

main 215 branches 17 tags

Switch branches/tags

Find a tag

Branches Tags

v1.12.0

v1.11.0

v1.11.0-dbaas

v1.10.0

Generic	EKS	GKE
		gcloud auth login
	eksctl create cluster -f ~/cluster.yaml	gcloud container clusters create m
		gcloud container clusters get-crede
		kubectl create clusterrolebinding cl
git clone -b v1.12.0 https://github.com/percona/percona-xtradb-cluster-operator		
cd percona-xtradb-cluster-operator		
kubectl create namespace pxc		
kubectl config set-context \$(kubectl config current-context) --namespace=pxc		
kubectl apply -f deploy/crd.yaml	kubectl apply -f deploy/bundle.yaml	
kubectl apply -f deploy/rbac.yaml		
kubectl apply -f deploy/operator.yaml		
kubectl create -f deploy/secrets.yaml		
kubectl apply -f deploy/cr.yaml		

# Deploying on GKE

<https://docs.percona.com/percona-operator-for-mysql/pxc/gke.html>

```
git clone -b v1.12.0 https://github.com/percona/percona-xtradb-cluster-operator
cd percona-xtradb-cluster-operator

gcloud auth login

gcloud container clusters create nando-1 --project pl2023-k8s-tutorial --zone us-central1-a --cluster-version 1.23
--machine-type n1-standard-4 --num-nodes=3

gcloud container clusters get-credentials nando-1 --zone us-central1-a --project pl2023-k8s-tutorial

kubectl create clusterrolebinding cluster-admin-binding --clusterrole cluster-admin --user $(gcloud config
get-value core/account)

kubectl create namespace pxc

kubectl config set-context $(kubectl config current-context) --namespace=pxc

kubectl apply -f deploy/bundle.yaml

kubectl apply -f deploy/cr.yaml
```

\* Remember to adjust: **cluster name**, **project name**, and possibly **zone** `kubectl get pods -w`



## Accessing test VMs

---

<https://bit.ly/45jtWwX>

```
chmod +x 600 p12023
```

```
ssh -i p12023 ubuntu@<IP>
```

# Deploying on GKE

```
$ gcloud container clusters list
```

NAME	LOCATION	MASTER_VERSION	MASTER_IP	MACHINE_TYPE	NODE_VERSION	NUM_NODES	STATUS
nando-1	us-central1-a	1.23.17-gke.300	34.122.11.227	n1-standard-4	1.23.17-gke.300	3	RUNNING

N1 standard

N1 high-memory

N1 high-cpu

N1 shared-core

N1 standard machine types have 3.75 GB of system memory per vCPU.

Machine types	vCPUs*	Memory (GB)	Max number of persistent disks (PDs) <sup>†</sup>	Max total PD size (TB)	Local SSD	Maximum egress bandwidth (Gbps) <sup>‡</sup>	Tier 1 egress bandwidth (Gbps)
n1-standard-1	1	3.75	128	257	Yes	2	N/A
n1-standard-2	2	7.50	128	257	Yes	10	N/A
n1-standard-4	4	15	128	257	Yes	10	N/A

# Deploying on GKE

<https://console.cloud.google.com>

Google Cloud consultants

Cloud overview >

View all products

PINNED

Pin your top products here

MORE PRODUCTS ^

COMPUTE

Compute Engine >

Kubernetes Engine >

VMware Engine

Clusters

Workloads

Services & Ingress

Applications

Secrets & ConfigMaps

Storage

Object Browser

Migrate to Containers

Backup for GKE **NEW**

Config Management

Security Posture

Kubernetes Engine

Kubernetes cl...

REFRESH

OPERATIONS

HELP ASSISTANT

LEARN

Clusters

Workloads

Services & Ingress

Applications

Secrets & ConfigMaps

Storage

Object Browser

OVERVIEW

OBSERVABILITY

COST OPTIMIZATION

Filter Enter property name or value

Status	Name	Location	Number of nodes	Total vCPUs	Total memory
<input type="checkbox"/>	training	us-central1-a	3	12	45 GB
<input type="checkbox"/>	nando-1	us-central1-a	3	12	45 GB

nando-1

DETAILS

NODES

STORAGE

LOGS

Node Pools

Filter Filter node pools

Name	Status	Version	Number of nodes	Machine type	Image type	Autoscaling	Pod IP address range
gke-nando-1-default-pool-sf8ba7ba-7wdd	Ok	1.23.12-gke.100	3	n1-standard-4	Container-Optimized OS with containerd (cos_containerd)	Off	10.104.0.0/14

Nodes

Filter Filter nodes

Name	Status	CPU requested	CPU allocatable	Memory requested	Memory allocatable	Storage
gke-nando-1-default-pool-sf8ba7ba-7wdd	Ready	2.1 CPU	3.92 CPU	2.7 GB	12.97 GB	
gke-nando-1-default-pool-sf8ba7ba-7wdd	Ready	2.11 CPU	3.92 CPU	2.69 GB	12.97 GB	
gke-nando-1-default-pool-sf8ba7ba-7wdd	Ready	1.99 CPU	3.92 CPU	2.71 GB	12.97 GB	

# Deploying on GKE

<input type="checkbox"/> Status	Name ↑	Location	Number of nodes	Total vCPUs	Total memory
<input type="checkbox"/>	<a href="#">nando-1</a>	us-central1-a	3	12	45 GB

## Nodes

```
$ kubectl get nodes
```


Filter Filter nodes

Name ↑

[gke-nando-1-default-pool-81437ff3-6j1f](#)

[gke-nando-1-default-pool-81437ff3-7w1x](#)

[gke-nando-1-default-pool-81437ff3-nsw1](#)



# Deploying on GKE

```
$ kubectl get pxc
```

NAME	ENDPOINT	STATUS	PXC	PROXYSQL	HAPROXY	AGE
cluster1	cluster1-haproxy.pxc	ready	3		3	7m

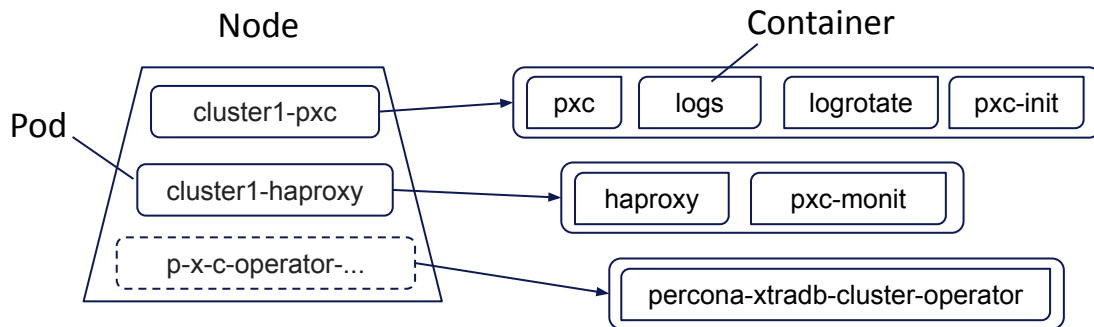
```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
cluster1-haproxy-0	2/2	Running	0	6m1s
cluster1-haproxy-1	2/2	Running	0	4m44s
cluster1-haproxy-2	2/2	Running	0	4m18s
cluster1-pxc-0	3/3	Running	0	6m1s
cluster1-pxc-1	3/3	Running	0	4m48s
cluster1-pxc-2	3/3	Running	0	3m36s
percona-xtradb-cluster-operator-5dbc998f8b-mm7wx	1/1	Running	0	6m36s

# Deploying on GKE

```
$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
cluster1-haproxy-0	2/2	Running	0	9m1s	10.104.1.7	gke-nando-1-default-pool-81437ff3-6j1f	<none>	<none>
cluster1-haproxy-1	2/2	Running	0	7m44s	10.104.0.7	gke-nando-1-default-pool-81437ff3-7w1x	<none>	<none>
cluster1-haproxy-2	2/2	Running	0	7m18s	10.104.2.5	gke-nando-1-default-pool-81437ff3-nsw1	<none>	<none>
cluster1-pxc-0	3/3	Running	0	9m1s	10.104.2.4	gke-nando-1-default-pool-81437ff3-nsw1	<none>	<none>
cluster1-pxc-1	3/3	Running	0	7m48s	10.104.1.8	gke-nando-1-default-pool-81437ff3-6j1f	<none>	<none>
cluster1-pxc-2	3/3	Running	0	6m36s	10.104.0.8	gke-nando-1-default-pool-81437ff3-7w1x	<none>	<none>
percona-xtradb-cluster-operator-5dbc998f8b-mm7wx	1/1	Running	0	9m36s	10.104.1.6	gke-nando-1-default-pool-81437ff3-6j1f	<none>	<none>



# Affinity

<https://docs.percona.com/percona-operator-for-mysql/pxc/constraints.html>

The *mysql* section of the *cr.yaml* file allows for the customization of node affinity:

- **affinity** rules make a pod eligible to run on a node hosting other pods labelled a certain way
- **anti-affinity** is the opposite: define the criteria that renders a pod ineligible to run on a given node
  - the standard employed in Percona operators

```
pxc:
  (...)
  affinity:
    antiAffinityTopologyKey: "kubernetes.io/hostname"
```



\* It is also possible to use standard Kubernetes constraints to define affinity rules, a more advanced approach.

Pods will avoid residing within the same **host**

# K9s (1)

k9s is a terminal client tool to ease the management of K8s:

<https://github.com/derailed/k9s>

```
Context: gke_consultants-206215_us-central1-a_nando-1
Cluster: gke_consultants-206215_us-central1-a_nando-1
User:    gke_consultants-206215_us-central1-a_nando-1
K9s Rev: v0.26.3 ⚡v0.26.7
K8s Rev: v1.23.13-gke.900
CPU:     4%
MEM:     9%
```

```
<0> all      <a> Attach
<1> pxc      <ctrl-d> Delete
<2> default  <d> Describe
           <e> Edit
           <?> Help
           <ctrl-k> Kill
```



## Pods(pxc)[7]

NAME ↑	PF	READY	RESTARTS	STATUS	CPU	MEM	%CPU/R	%CPU/L	%MEM/R	%MEM/L	IP
cluster1-haproxy-0	●	2/2	0	Running	55	8	9	n/a	0	n/a	10.1
cluster1-haproxy-1	●	2/2	0	Running	57	8	9	n/a	0	n/a	10.1
cluster1-haproxy-2	●	2/2	0	Running	50	8	8	n/a	0	n/a	10.1
cluster1-pxc-0	●	3/3	0	Running	21	447	2	n/a	39	n/a	10.1
cluster1-pxc-1	●	3/3	0	Running	21	437	2	n/a	38	n/a	10.1
cluster1-pxc-2	●	3/3	0	Running	18	442	1	n/a	38	n/a	10.1
percona-xtradb-cluster-operator-5dbc998f8b-mm7wx	●	1/1	0	Running	11	22	11	5	113	4	10.1

<pod>



# K9s (2)

## Containers in a pod:

Context: gke\_consultants-206215\_us-central1-a\_nando-1  
Cluster: gke\_consultants-206215\_us-central1-a\_nando-1  
User: gke\_consultants-206215\_us-central1-a\_nando-1  
K9s Rev: v0.26.3 ⚡ v0.26.7  
K8s Rev: v1.23.13-gke.900  
CPU: 5%  
MEM: 9%

<a> Attach  
<?> Help  
<l> Logs  
<p> Logs Previous  
<shift-f> PortForward  
<s> Shell

<f> Show...



Containers(pxc/cluster1-pxc-0)[4]

NAME ↑	PF	IMAGE	READY	STATE	INIT	RESTARTS	PROBES(L:
logrotate	●	percona/percona-xtradb-cluster-operator:1.11.0-logcollector	true	Running	false	0	off:off
logs	●	percona/percona-xtradb-cluster-operator:1.11.0-logcollector	true	Running	false	0	off:off
pxc	●	percona/percona-xtradb-cluster:8.0.27-18.1	true	Running	false	0	on:on
pxc-init	●	percona/percona-xtradb-cluster-operator:1.11.0	true	Completed	true	0	off:off

<pod>

<containers>

# K9s (3)

---

Some interesting shortcuts:

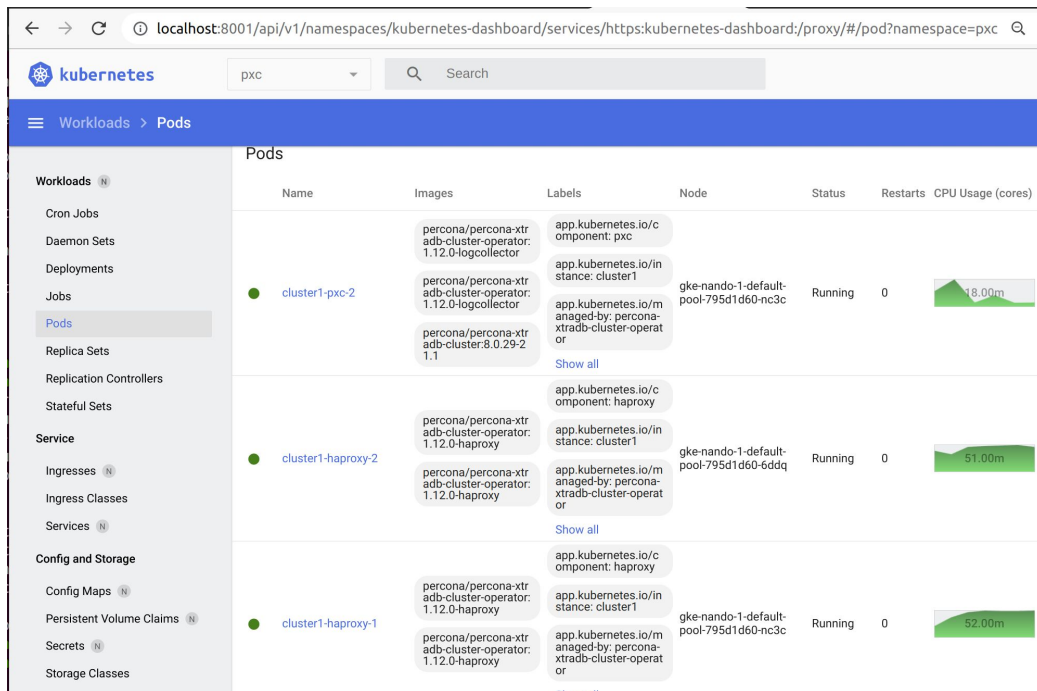
- :node → show nodes
  - :namespace → show namespaces
  - :pod → show pods
  - :service → show services
- 
- Can edit configuration
  - Can access logs
  - Can open shell in container
  - Can redirect TCP ports

# Kubernetes Dashboard

A "(...) general-purpose, web-based UI for Kubernetes clusters.

- <https://github.com/kubernetes/dashboard>

\* Requires the creation of a special Service Account granted with *ClusterRoleBinding* privilege.



The screenshot displays the Kubernetes Dashboard interface. The top navigation bar shows the 'Workloads' section selected, with 'Pods' as the active sub-section. The left sidebar contains a menu with categories: Workloads (Cron Jobs, Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, Stateful Sets), Service (Ingresses, Ingress Classes, Services), and Config and Storage (Config Maps, Persistent Volume Claims, Secrets, Storage Classes). The main content area lists Pods with columns for Name, Images, Labels, Node, Status, Restarts, and CPU Usage (cores). Three pods are visible, all in a 'Running' state. Each pod entry includes a green status dot, a pod name (e.g., cluster1-pxc-2), a list of images (percona/percona-xtradb-cluster-operator, 1.12.0-logcollector), labels (app.kubernetes.io/component: pxc, app.kubernetes.io/instance: cluster1, app.kubernetes.io/managed-by: percona-xtradb-cluster-operator), the node name (gke-nando-1-default-pool-795d1d60-nc3c), and a CPU usage graph showing 8.00m.

Name	Images	Labels	Node	Status	Restarts	CPU Usage (cores)
cluster1-pxc-2	percona/percona-xtradb-cluster-operator: 1.12.0-logcollector	app.kubernetes.io/component: pxc app.kubernetes.io/instance: cluster1 app.kubernetes.io/managed-by: percona-xtradb-cluster-operator	gke-nando-1-default-pool-795d1d60-nc3c	Running	0	8.00m
cluster1-haproxy-2	percona/percona-xtradb-cluster-operator: 1.12.0-haproxy	app.kubernetes.io/component: haproxy app.kubernetes.io/instance: cluster1 app.kubernetes.io/managed-by: percona-xtradb-cluster-operator	gke-nando-1-default-pool-795d1d60-6ddq	Running	0	51.00m
cluster1-haproxy-1	percona/percona-xtradb-cluster-operator: 1.12.0-haproxy	app.kubernetes.io/component: haproxy app.kubernetes.io/instance: cluster1 app.kubernetes.io/managed-by: percona-xtradb-cluster-operator	gke-nando-1-default-pool-795d1d60-nc3c	Running	0	52.00m

---

# Accessing the database

# Secrets

<https://docs.percona.com/percona-operator-for-mysql/pxc/users.html>

```
$ kubectl get secrets
```

NAME	TYPE	DATA
AGE		
cluster1-secrets	Opaque	7
8d		
cluster1-ssl	kubernetes.io/tls	3
8d		
cluster1-ssl-internal	kubernetes.io/tls	3
8d		

```
$ kubectl get secrets cluster1-secrets -o yaml
```

```
apiVersion: v1
data:
  clustercheck: elNva09De1VQMmU3TURTVg==
  percona-xtradb-cluster-operator-token: x2qc9
  monitor: R3IIMk5bKsyPMQx33Zkfr1StOA=
  operator: OE1USkVfZj9HPWNlW2xrKGZi
  proxyadmin: NHlAPUFpQWpNbHBHWlZDVjx9Pw==
  replication: MU9qbKJaSkNpRUZWbUY3aA==
  root: WW1wVVZpVERZN1RKc1lCaXJXCg==
  xtrabackup: S2xHQ1NNOEVyNlNwU1BXMkY=
  (...)

```

First, retrieve the encoded password for the MySQL *root* user

```
$ echo "WW1wVVZpVERZN1RKc1lCaXJXCg==" | base64
-d
YmpUViTDY7TJsYBirW
```

# Accessing the database from within K8s

---

The “unconventional” way for K8s: directly to one of the nodes

```
$ kubectl exec -it cluster1-pxc-1 -c pxc -- bash

bash-4.4$ mysql -uroot -pYmpUViTdY7TJsYBirW
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 7915
Server version: 8.0.27-18.1 Percona XtraDB Cluster (GPL), Release rel18, Revision ac35177, WSREP
version 26.4.3
...
```

# Accessing the database from within K8s

---

The “unconventional” way for K8s: directly to one of the nodes

*one-liner*

```
$ kubectl exec -it cluster1-pxc-1 -c pxc -- mysql -uroot -pYmpUViTdY7TJsYBirW
```

# Accessing the database from within K8s

<https://docs.percona.com/percona-operator-for-mysql/pxc/haproxy-conf.html>

## Connecting to the cluster through HAproxy

```
$ kubectl get service cluster1-haproxy
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
cluster1-haproxy	ClusterIP	10.108.6.176	<none>	3306/TCP,3309/TCP,33062/TCP,33060/TCP

Annotations: primary (pointing to 3306/TCP), proxy protocol (pointing to 3309/TCP)

```
$ kubectl get service cluster1-haproxy-replicas
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
cluster1-haproxy-replicas	ClusterIP	10.108.3.134	<none>	3306/TCP	104m

Annotation: replicas (pointing to 3306/TCP)

\$ kubectl get services



# Accessing the database from within K8s

## Connecting to the cluster through HAproxy

Docker image with the *mysql* 8.0 command-line client

```
$ kubectl run -i --rm --tty percona-client --image= percona:8.0 --restart=Never -- bash -il
```

```
$ mysql -h cluster1-haproxy -uroot -pYmpUViTdY7TJsYBirW -P3306 -e 'select @@hostname'
+-----+
| @@hostname |
+-----+
| cluster1-pxc-0 |
+-----+
```

round-robin

```
$ mysql -h cluster1-haproxy-replicas -uroot -pYmpUViTdY7TJsYBirW -P3306 -e 'select @@hostname'
+-----+
| @@hostname |
+-----+
| cluster1-pxc-1 |
+-----+
```

# Access from your notebook: port-forwarding

---

```
$ kubectl port-forward svc/cluster1-haproxy 8080:3306
```

```
Forwarding from 127.0.0.1:8080 -> 3306
```

```
Forwarding from [::1]:8080 -> 3306
```

```
$ mysql -h 127.0.0.1 -P 8080 -uroot -pYmpUViTDY7TJsYBirW
```

```
mysql: [Warning] Using a password on the command line interface can be insecure.
```

```
Welcome to the MySQL monitor.  Commands end with ; or \g.
```

```
Your MySQL connection id is 7552
```

```
Server version: 8.0.29-21.1 Percona XtraDB Cluster (GPL), Release rel21, Revision 250bc93, WSREP  
version 26.4.3
```

---

# Making customizations

# Access from outside K8s: Expose


- Uses the K8s load balancer utility
- Allows external access

But the right thing to do is to access the database through **HProxy\***:

Edit `deploy/cr.yaml`,  
uncommenting these 3 lines:

```
pxc:
  (...)
  expose:
    enabled: true
    type:
      LoadBalancer
```

```
haproxy:
  (...)
  serviceType: LoadBalancer
  externalTrafficPolicy: Cluster
  serviceAnnotations:
    networking.gke.io/load-balancer-type:
      "Internal"
  (...)
  replicasServiceType: LoadBalancer
  replicasExternalTrafficPolicy: Cluster
  replicasServiceAnnotations:
    networking.gke.io/load-balancer-type:
      "Internal"
```



```
$ kubectl apply -f deploy/cr.yaml
```

\* or ProxySQL, if that's what  
you chose instead.

# Access from outside K8s: Expose

- Just for MySQL access, doesn't go through HAProxy or ProxySQL

```
$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
cluster1-haproxy	ClusterIP	10.108.6.176	10.128.0.118	3306/TCP, 3309/TCP, 33062/TCP, 33060/TCP	147m
cluster1-haproxy-replicas	ClusterIP	10.108.3.134	10.128.0.119	3306/TCP	147m
cluster1-pxc	ClusterIP	None	<none>	3306/TCP, 33062/TCP, 33060/TCP	147m
cluster1-pxc-0	LoadBalancer	10.108.0.55	34.122.165.169	3306:32650/TCP	7m10s
cluster1-pxc-1	LoadBalancer	10.108.10.192	35.202.255.108	3306:30203/TCP	7m9s
cluster1-pxc-2	LoadBalancer	10.108.6.6	34.170.91.92	3306:32272/TCP	7m9s
cluster1-pxc-unready	ClusterIP	None	<none>	3306/TCP, 33062/TCP, 33060/TCP	147m
percona-xtradb-cluster-operator	ClusterIP	10.108.7.100	<none>	443/TCP	147m

Get more details about a specific service:

```
$ kubectl get services cluster1-pxc-0 -o wide
```

```
$ kubectl get services cluster1-pxc-0 -o yaml
```

# Customizing MySQL

<https://docs.percona.com/percona-operator-for-mysql/pxc/options.html>

```
spec:
  pxc:
    configuration: |
      [mysqld]
#      wsrep_debug=CLIENT
#      wsrep_provider_options="gcache.size=1G; gcache.recover=yes"
      innodb_buffer_pool_size=2G
      tmp_table_size=32M
```



```
$ kubectl apply -f deploy/cr.yaml
```

Complete reference for the Custom Resource (CR) options:

<https://docs.percona.com/percona-operator-for-mysql/pxc/operator.html>

# Upgrading MySQL root password


Upgrading MySQL *root* password:

not on macOS

OR:

```
$ echo -n "au78yEJEKGs6R96b" | base64  
--wrap=0  
YXU3OHlFskVLR3M2Ujk2Yg==
```

```
$ kubectl edit
```



```
$ kubectl patch secret/cluster1-secrets -p '{"data":{"root": " YXU3OHlFskVLR3M2Ujk2Yg=="}}'
```

OR:

*kubectl apply -f update\_secret.yaml*

```
apiVersion: v1  
kind: Secret  
metadata:  
  name: cluster1-secrets  
type: Opaque  
data:  
  root: YXU3OHlFskVLR3M2Ujk2Yg==
```

\* Wait a few seconds/minutes to take effect!

```
$ kubectl exec -it cluster1-pxc-0 -c  
pxc -- mysql -uroot -p au78yEJEKGs6R96b
```

# Upgrading MySQL root password

An alternative way:

```
$ cat deploy/secrets2.yaml
apiVersion: v1
kind: Secret
metadata:
  name: cluster1-secrets
type: Opaque
stringData:
  root: YmpUViTdY7TJsYBirW
```

data  
vs  
stringData

```
$ kubectl apply -f deploy/secrets2.yaml
```



# Cluster management

# Cluster management

---

- Backup and restore
- Horizontal and vertical scaling
- Adding sidecar containers
- Pausing a cluster
- Deleting a cluster

# Backup

<https://docs.percona.com/percona-operator-for-mysql/pxc/backups.html>

```
backup:
  image:
percona/percona-xtradb-cluster-operator:1.12.0-pxc8.0-backup
(...)
  pitr:
    enabled: false
    storageName: STORAGE-NAME-HERE
    timeBetweenUploads: 60
(...)
  storages:
    s3-us-west:
      type: s3
      verifyTLS: true
(...)
  s3:
    bucket: S3-BACKUP-BUCKET-NAME-HERE
    credentialsSecret: my-cluster-name-backup-s3
    region: us-west-2
  fs-pvc:
    type: filesystem
```

```
(...)
  schedule:
    - name: "sat-night-backup"
      schedule: "0 0 * * 6"
      keep: 3
      storageName: s3-us-west
    - name: "daily-backup"
      schedule: "0 0 * * *"
      keep: 5
      storageName: fs-pvc
```

# Backup

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
cluster1-haproxy-0	2/2	Running	0	3d
cluster1-haproxy-1	2/2	Running	0	3d
cluster1-haproxy-2	2/2	Running	0	3d
cluster1-pxc-0	3/3	Running	0	3d
cluster1-pxc-1	3/3	Running	0	3d
cluster1-pxc-2	3/3	Running	0	3d
percona-client	0/1	Completed	0	3d
percona-xtradb-cluster-operator-5dbc998f8b-mm7wx	1/1	Running	0	3d
xb-cron-cluster1-fs-pvc-20221126000-372f8-gzsqx	0/1	Completed	0	2d18h
xb-cron-cluster1-fs-pvc-20221127000-372f8-152qm	0/1	Completed	0	42h
xb-cron-cluster1-fs-pvc-20221128000-372f8-rww6g	0/1	Completed	0	18h
xb-cron-cluster1-s3-us-west-20221126000-3d2dv-9ttzw	0/1	CreateContainerConfigError	0	2d18h

```
$ kubectl get pxc-backups
```

NAME	CLUSTER	STORAGE	DESTINATION	STATUS	COMPLETED
cron-cluster1-fs-pvc-20221126000-372f8 2d17h	cluster1	fs-pvc	pvc/xb-cron-cluster1-fs-pvc-20221126000-372f8	Succeeded	2d17h
cron-cluster1-fs-pvc-20221127000-372f8 41h	cluster1	fs-pvc	pvc/xb-cron-cluster1-fs-pvc-20221127000-372f8	Succeeded	41h
cron-cluster1-fs-pvc-20221128000-372f8	cluster1	fs-pvc	pvc/xb-cron-cluster1-fs-pvc-20221128000-372f8	Succeeded	17h

# Backup

1) Set-up credentials to S3 (*kubectl apply -f* deploy/backup-s3.yaml):

```
apiVersion: v1
kind: Secret
metadata:
  name: cluster1-s3-credentials
type: Opaque
data:
  AWS_ACCESS_KEY_ID: <base64 encoded>
  AWS_SECRET_ACCESS_KEY: <base64 encoded>
```

3) Create a manual backup (*kubectl apply -f* deploy/backup/backup.yaml):

```
apiVersion: pxc.percona.com/v1
kind: PerconaXtraDBClusterBackup
metadata:
  name: backup1
# finalizers:
#   - delete-backup
spec:
  pxcCluster: cluster1
  storageName: s3-us-west
```

2) Configure backup and storage on (*kubectl apply -f* deploy/cr.yaml):

```
backup:
  enabled: true
  (...)
  storages:
    s3-us-west:
      type: s3
      verifyTLS: true
  (...)
s3:
  bucket: nando-s3
  credentialsSecret: cluster1-s3-credentials
  region: us-east-1
#   endpointUrl: https://...
```

```
$ kubectl get pxc-backup
NAME          STORAGE    DESTINATION
STATE         COMPLETED AGE
backup1       s3-us-west
s3://nando-s3/cluster1-2023-04-18-20:08:20-full  Succeeded
27m          27m
```

# Restore

Restoring from a backup previously taken **OR** Restoring from a backup from a different source

Configure restore (*kubectl apply -f* deploy/backup/restore.yaml):

```
apiVersion: pxc.percona.com/v1
kind: PerconaXtraDBClusterRestore
metadata:
  name: restore1
spec:
  pxcCluster: cluster1
  backupName: backup1
```

```
apiVersion: pxc.percona.com/v1
kind: PerconaXtraDBClusterRestore
metadata:
  name: restore1
spec:
  pxcCluster: cluster1
  backupSource:
    destination:
      s3://nando-s3/cluster1-2023-04-19-16:52:35-full
    s3:
      bucket: nando-s3
      credentialsSecret: cluster1-s3-credentials
      region: us-east-1
```

```
$ kubectl get pxc-restore
NAME          CLUSTER   STATUS          COMPLETED
AGE
restore1     cluster1  Stopping Cluster
7s
restore1     cluster1  Restoring
```

# Backup - binlogs

## PITR: binlog streaming only works with S3-like storage

*kubectl apply -f deploy/cr.yaml*

```
backup:
  image: percona/percona-xtradb-cluster-operator:1.12.0-pxc8.0-backup
#   backoffLimit: 6
#   serviceAccountName: percona-xtradb-cluster-operator
#   imagePullSecrets:
#     - name: private-registry-credentials
pitr:
  enabled: true
  storageName: s3-us-west
  timeBetweenUploads: 60
```

\$ kubectl get pods

NAME	READY	STATUS	RESTARTS	AGE
cluster1-haproxy-0	2/2	Running	0	19m
cluster1-haproxy-1	2/2	Running	0	18m
cluster1-haproxy-2	2/2	Running	0	18m
cluster1-pitr-76b97bd96f-vzbzx	1/1	Running	0	2m2s
cluster1-pxc-0	3/3	Running	0	19m
cluster1-pxc-1	3/3	Running	0	19m
cluster1-pxc-2	3/3	Running	0	17m
percona-xtradb-cluster-operator-77bf8b9df5-kgxr6	1/1	Running	0	106m
restore-iob-restore1-cluster1-zd8dz	0/1	Completed	0	20m

# Restore - PITR

*kubectl apply -f deploy/backup/restore.yaml*

```
apiVersion: pxc.percona.com/v1
kind: PerconaXtraDBClusterRestore
metadata:
  name: restore2
spec:
  pxcCluster: cluster1
  backupSource:
    destination: s3://nando-s3/cluster1-2023-04-19-16:52:35-full
    s3:
      bucket: nando-s3
      credentialsSecret: cluster1-s3-credentials
      region: us-east-1
  pitr:
    type: transaction
    gtid: "81386cf6-de2c-11ed-972e-5f2edd7af261:7"
    backupSource:
      storageName: s3-us-west
```

```
$ kubectl get pods
NAME
cluster1-haproxy-0
cluster1-haproxy-1
cluster1-haproxy-2
cluster1-pxc-0
cluster1-pxc-1
cluster1-pxc-2
percona-xtradb-cluster-operator-77bf8b9df5-6xkl7
pitr-job-restore2-cluster1-4mwx7
restore-job-restore2-cluster1-9754h
```

```
$ aws s3 ls s3://nando-s3/
PRE cluster1-2023-04-19-16:52:35-full.sst_info/
PRE cluster1-2023-04-19-16:52:35-full /
(...)
2023-04-19 13:52:59      25909 cluster1-2023-04-19-16:52:35-full.md5
2023-04-19 13:55:37       38 last-binlog-set-81386cf6-de2c-11ed-972e-5f2edd7af261
```



# Restore - retrieving the backup file

The operator includes a script that facilitates copying over backups to a local computer:

```
$ kubectl get pxc-backup
```

NAME	CLUSTER	STORAGE	DESTINATION	STATUS	COMPLETED	AGE
<b>cron-cluster1-fs-pvc-20234200040-372f8</b>	cluster1	fs-pvc	pvc/xb-cron-cluster1-fs-pvc-20234200040-372f8	Succeeded	17h	17h

```
$ deploy/backup/copy-backup.sh cron-cluster1-fs-pvc-20234200040-372f8 /tmp/backup
```

```
Log: /tmp/tmp.tCcrpi1R1G/log
```

```
pvc/xb-cron-cluster1-fs-pvc-20234200040-372f8pod "backup-access" deleted
```

```
pod/backup-access created
```

```
Starting pod..[done]
```

```
Downloading started
```

```
tar: Removing leading `/' from member names
```

```
Downloading finished
```

```
pod "backup-access" deleted
```

You can recover data locally with following commands:

```
$ service mysqld stop
```

```
$ rm -rf /var/lib/mysql/*
```

```
$ cat /tmp/backup/xtrabackup.stream | xbstream --decompress -x -C /var/lib/mysql
```

```
$ xtrabackup --prepare --target-dir=/var/lib/mysql
```

```
$ chown -R mysql:mysql /var/lib/mysql
```

```
$ service mysqld start
```

```
$ ls /tmp/backup
```

```
lost+found md5sum.txt
```

```
sst_info xtrabackup.stream
```

# Restoring the backup in a sidecar container

It's possible to prepare a backup that was taken locally and restore it on a separate container (you may want to do a copy of the pvc hosting the backup first!):

```
apiVersion: v1
kind: Pod
metadata:
  name: backup-access
spec:
  containers:
  - name: xtrabackupview
    image: percona/percona-xtrabackup:8.0.29
    command: ["/bin/sh"]
    args: ["-c", "while true; do trap 'exit 0' SIGINT SIGTERM SIGQUIT SIGKILL; done;"]
    volumeMounts:
    - name: backup
      mountPath: /backup
  restartPolicy: Never
  volumes:
  - name: backup
    persistentVolumeClaim:
      claimName: pvc-c8420f88-f74b-4cc8-b514-a12b54c0e7d4
```

*\$ kubectl apply -f backup-access.yaml*

# Restoring the backup in a sidecar container

---

Decompress and prepare the backup:

```
$ kubectl exec -it backup-access -- sh  
  
sh-4.4# cd /backup  
sh-4.4# xbstream -x < xtrabackup.stream  
sh-4.4# xtrabackup --prepare --target-dir=/backup  
sh-4.4# chown 1001:1001 -R /backup
```

You can now delete that container - the PVC will remain intact:

```
$ kubectl delete pod backup-access
```

# Restoring the backup in a sidecar container

---

Edit the container definition file (*backup-access.yaml*) as to change the image used from *xtrabackup* to a Percona Server one:

```
spec.image: percona/percona-server:8.0.29
```

Deploy the container again:

```
$ kubectl apply -f backup-access.yaml
```

Get inside it, edit the *my.cnf* sample file included, commenting all variables that are PXC-specific:

```
$ kubectl exec -it backup-access -- sh
sh-4.4# vi /backup/backup-my.cnf
```

and, finally, start a *mysqld* instance with the datadir:

```
$ mysqld --defaults-file=/backup/backup-my.cnf --datadir=/backup
```

# Increasing memory, CPU, and storage

<https://docs.percona.com/percona-operator-for-mysql/pxc/scaling.html>

Filter nando-1 Filter nodes × ?

Name ↑	CPU requested	CPU allocatable	Memory requested	Memory allocatable
<a href="#">gke-nando-1-default-pool-81437ff3-6j1f</a>	2 CPU	3.92 CPU	2.71 GB	12.97 GB
<a href="#">gke-nando-1-default-pool-81437ff3-7w1x</a>	2.11 CPU	3.92 CPU	2.71 GB	12.97 GB
<a href="#">gke-nando-1-default-pool-81437ff3-nsw1</a>	2.09 CPU	3.92 CPU	2.68 GB	12.97 GB

Filter nando-1 Filter persistent volume claims

<input type="checkbox"/>	Name ↑	Phase	Volume	Storage class	Namespace
<input type="checkbox"/>	<a href="#">datadir-cluster1-pxc-0</a>	✔ Bound	<a href="#">pvc-13818b82-b436-4869-a8eb-d7ea4a79728a</a>	<a href="#">standard</a>	pxc
<input type="checkbox"/>	<a href="#">datadir-cluster1-pxc-1</a>	✔ Bound	<a href="#">pvc-9b8b2e8b-fa16-404a-95ef-6ca4ae549c44</a>	<a href="#">standard</a>	pxc
<input type="checkbox"/>	<a href="#">datadir-cluster1-pxc-2</a>	✔ Bound	<a href="#">pvc-3804116c-5389-4cb8-bbf3-ca7d6c8133ce</a>	<a href="#">standard</a>	pxc

# Increasing memory, CPU, and storage

```
$ kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
datadir-cluster1-pxc-0	Bound	pvc-13818b82-b436-4869-a8eb-d7ea4a79728a	6Gi	RWO	standard	17h
datadir-cluster1-pxc-1	Bound	pvc-9b8b2e8b-fa16-404a-95ef-6ca4ae549c44	6Gi	RWO	standard	16h
datadir-cluster1-pxc-2	Bound	pvc-3804116c-5389-4cb8-bbf3-ca7d6c8133ce	6Gi	RWO	standard	16h

```
$ kubectl exec -it cluster1-pxc-0 -c pxc -- df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
overlay	95G	6.0G	89G	7%	/
tmpfs	64M	0	64M	0%	/dev
tmpfs	7.4G	0	7.4G	0%	/sys/fs/cgroup
/dev/sda1	95G	6.0G	89G	7%	/tmp
shm	64M	0	64M	0%	/dev/shm
/dev/sdb	5.9G	325M	5.5G	6%	/var/lib/mysql
tmpfs	13G	0	13G	0%	/etc/mysql/vault-keyring-secret
tmpfs	13G	32K	13G	1%	/etc/mysql/mysql-users-secret
tmpfs	13G	12K	13G	1%	/etc/mysql/ssl-internal
tmpfs	13G	12K	13G	1%	/etc/mysql/ssl
tmpfs	13G	12K	13G	1%	/run/secrets/kubernetes.io/serviceaccount
tmpfs	7.4G	0	7.4G	0%	/proc/acpi
tmpfs	7.4G	0	7.4G	0%	/proc/scsi
tmpfs	7.4G	0	7.4G	0%	/sys/firmware

# Increasing memory, CPU, and storage

```
$ cat deploy/cr.yaml
(...)
pxc:
(...)
  resources:
    requests:
      memory: 1G      -> 2G
      cpu: 600m       -> 1
#    ephemeral-storage: 1G
#  limits:
#    memory: 1G
(...)
  volumeSpec:
(...)
    persistentVolumeClaim:
(...)
      resources:
        requests:
          storage: 6G
```

*not enough!*

-> 10G

```
$ kubectl apply -f
deploy/cr.yaml
```

**NOTE:** On **AWS** *allowVolumeExpansion* is set to False by default

```
$ kubectl get storageclass
```

```
$ kubectl edit storageclass gp2
```

```
allowVolumeExpansion: true
```

```
$ kubectl edit pvc datadir-cluster1-pxc-2
```

```
spec:
  accessModes:
  - ReadWriteOnce
  resources:
    requests:
      storage: 10G
```

\* The datadir PVs should be expanded to the new size automatically within a few minutes. If not, you can always perform a rolling restart.

# Scaling the cluster up

1) First, you may need to increase the “number” of nodes in the K8s cluster:

```
$ gcloud container clusters resize nando-1 --project consultants-206215 --zone us-central1-a --num-nodes=5
```

2) Second, choose one of 3 ways to increase the “size” of the cluster:

a)

```
$ cat deploy/cr.yaml
(...)
pxc:
(...)
  size: 5
```

```
$ kubectl apply -f
deploy/cr.yaml
```

b)

```
$ kubectl scale --replicas=5 pxc/cluster1
```

c)

```
$ kubectl edit pxc cluster1

spec:
  pxc:
    size: 5
```



# Scaling the cluster down

---

To scale down, revert the previous steps:

1) Decrease the “size” of the cluster:

```
$ kubectl scale --replicas=3 pxc/cluster1
```

2) Optionally, decrease the “number” of nodes in the K8s cluster:

```
$ gcloud container clusters resize nando-1 --project consultants-206215 --zone us-central1-a --num-nodes=3
```

# Sidecars

<https://docs.percona.com/percona-operator-for-mysql/pxc/sidecar.html>

- Sidecar container allow access to the datafiles of another container
- Can be useful for debugging, monitoring and performance tuning
- pmm-client runs as a sidecar
- Sidecars run as part of the pods

Adding a sidecar container (*kubectl apply -f* `deploy/cr.yaml`):

**Careful:** it does restart the pods!

```
spec:
  pxc:
    (...)
    sidecars:
    - name: sysbench
      image: perconalab/sysbench
      command: ["sleep", "30d"]
```

```
$ kubectl exec -it cluster1-pxc-2 -c sysbench -- bash
```

# Sidecars

---

## Running sysbench through HAproxy:

1) First, connect to the database and create the target database:

```
$ kubectl run -i --rm --tty percona-client1 --image=percona:8.0 --restart=Never -- bash -il
# mysql -h cluster1-haproxy -uroot -pYmpUViTdY7TJsYBirW -P3306 -e 'create database sbtest'
```

2) Then, on the *sysbench* session, create the schema and populate the data:

```
# LUA_PATH=/sysbench/sysbench-tpcc/?..lua sysbench-tpcc/tpcc.lua
--mysql-host=cluster1-haproxy --mysql-user=root --mysql-password=YmpUViTdY7TJsYBirW
--mysql-port=3306 --scale=10 --mysql-db=sbtest --db-driver=mysql --force-pk=1  prepare
```

# Sidecars

To make things easier, you may

add an environment variable to the sidecar container (deploy/cr.yaml):

*kubectl apply -f*

```
spec:
  pxc:
    (...)
    sidecars:
      - name: sysbench
        command: ["sleep", "30d"]
        env:
          - name: "LUA_PATH"
            value: "/sysbench/sysbench-tpcc/?.lua"
        image: perconalab/sysbench
```

3) Run the workload:

```
$ kubectl exec -it cluster1-pxc-2 -c sysbench -- sysbench-tpcc/tpcc.lua
--mysql-host=cluster1-haproxy --mysql-user=root --mysql-password=YmpUViTDY7TJsYBirW
--mysql-port=3306 --scale=10 --mysql-db=sbtest --db-driver=mysql --force-pk=1 run
```

# Sidecars

---

Or you can simply run Sysbench on a separate pod:

```
$ kubectl run sysbench1 --image=perconalab/sysbench --restart=Never  
--env="LUA_PATH=/sysbench/sysbench-tpcc/?.lua" --command --  
sysbench-tpcc/tpcc.lua --mysql-host=cluster1-haproxy --mysql-port=3306  
--mysql-user=root --mysql-password=YmpUViTDY7TJsYBirW --scale=10  
--mysql-db=sbtest --db-driver=mysql  
--force-pk=1 run
```

# Pausing the cluster

---

<https://docs.percona.com/percona-operator-for-mysql/pxc/pause.html>

Gracefully pause the cluster, for some maintenance or to “restart” it:

*kubectl apply -f*  
(deploy/cr.yaml):


```
spec:
  (...)
  allowUnsafeConfigurations: false
  pause: true
```

Setting it back to *false* starts the pods.

# Deleting

Deleting the PXC cluster:

```
$ kubectl delete -f deploy/cr.yaml
```



Unless the *finalizer* delete-pxc-pvc is used, data volumes won't be deleted

Deleting the operator:

```
$ kubectl delete -f deploy/operator.yaml
```

```
apiVersion: pxc.percona.com/v1
kind: PerconaXtraDBCluster
metadata:
  name: cluster1
  finalizers:
    - delete-pxc-pods-in-order
#    - delete-ssl
#    - delete-proxysql-pvc
#    - delete-pxc-pvc
```

Deleting the GKE cluster:

```
$ gcloud container clusters delete nando-1 --project consultants-206215 --zone us-central1-a
```

# Deleting

## Deleting the cluster

```
$ gcloud container clusters list
```

NAME	LOCATION	MASTER_VERSION	MASTER_IP	MACHINE_TYPE	NODE_VERSION	NUM_NODES	STATUS
training	us-central1-a	1.23.9-gke.900	34.72.11.250	n1-standard-4	1.23.9-gke.900	3	RUNNING
nando-1	us-central1-a	1.23.12-gke.100	34.121.252.81	n1-standard-4	1.23.12-gke.100	3	RUNNING

```
$ gcloud container clusters delete nando-1 --zone us-central1-a
```

The following clusters will be deleted.

- [nando-1] in [us-central1-a]

Do you want to continue (Y/n)? y

Deleting cluster nando-1...done.

Deleted

[<https://container.googleapis.com/v1/projects/consultants-206215/zones/us-central1-a/clusters/nando-1>].



Takes some time  
to complete!



# Troubleshooting

# Troubleshooting

---

- Crash recovery
- Debug and troubleshooting

# Troubleshooting

## Logs, logs, logs

```
$ kubectl logs cluster1-pxc-0 -c logs [ --previous ]
```

```
$ kubectl logs cluster1-pxc-0 -c  
pxc
```

```
$ kubectl logs  
percona-xtradb-cluster-operator-846bfc8f54-qfrhk
```

OR:

```
$ kubectl logs  
deploy/percona-xtradb-cluster-operator
```

## Describe:

```
$ kubectl describe pod  
cluster1-pxc-0
```

```
$ kubectl describe node  
cluster1-default-001-81427662-6i16
```

## Deployments (e.g.: changes applied to cr.yaml)

```
$ kubectl rollout status deployments percona-xtradb-cluster-operator
```

# Troubleshooting

## Percona tools: temporary *Percona-Toolkit* image available

```
$ kubectl run pt2 --image=perconalab/percona-server-mysql-operator:main-toolkit
--restart=Never --command sleep 20d
$ kubectl exec -it pt2 -- sh

sh-4.4$ pt-stalk --no-stalk --iterations=2 --sleep=30 --mysql-only --dest=/tmp/ptstalk --
--host=cluster1-haproxy --user=root --password=au78yEJEKGS6R96b
sh-4.4$ tar cvf /tmp/ptstalk-data.tar /tmp/ptstalk
sh-4.4$ exit

$ kubectl cp pt2:/tmp/ptstalk-data.tar ptstalk-data.tar
```

### pt-k8s-debug-collector

```
$ pt-k8s-debug-collector --cluster pxc.percona.com/cluster1 --namespace pxc
```



cluster-dump.tar.gz

# Troubleshooting

## Inspecting a K8s node:

```
$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE
READINESS GATES							
(...)							
cluster1-pxc-2	3/3	Running	0	145m	10.104.2.11	gke-nando-1-default-pool-e6acb66d-6lvz	<none>
(...)							<none>

```
$ kubectl debug node/ gke-nando-1-default-pool-e6acb66d-6lvz -it --image=ubuntu -- bash
```

```
Creating debugging pod node-debugger-gke-nando-1-default-pool-e6acb66d-6lvz-pbxc2 with container debugger on node  
gke-nando-1-default-pool-e6acb66d-6lvz.  
If you don't see a command prompt, try pressing enter.
```

```
root@gke-nando-1-default-pool-e6acb66d-6lvz:/# ps fax|grep "mysql\\haproxy"  
9909 ?      Ss      0:00 \_ haproxy -x /etc/haproxy/pxc/haproxy.sock -W -db -f /etc/haproxy/haproxy-global.cfg -f /etc/haproxy/pxc/haproxy.cfg -p  
/etc/haproxy/pxc/haproxy.pid -S /etc/haproxy/pxc/haproxy-main.sock  
13629 ?      Sl      0:11 | \_ haproxy -sf 486 -x /etc/haproxy/pxc/haproxy.sock -W -db -f /etc/haproxy/haproxy-global.cfg -f  
/etc/haproxy/pxc/haproxy.cfg -p /etc/haproxy/pxc/haproxy.pid -S /etc/haproxy/pxc/haproxy-main.sock  
11311 ?      Ssl     0:00 \_ go-cron 0 0 * * * sh -c logrotate -s /opt/percona/logrotate/logrotate.status  
/opt/percona/logrotate/logrotate-mysql.conf;/usr/bin/find /var/lib/mysql/ -name GRA_*.log -mtime +7 -delete  
11494 ?      Ssl     2:18 \_ mysqld  
231969 pts/0   S+      0:00 \_ grep --color=auto mysql\\haproxy
```

Sometimes it is possible to install additional software in this temporary, debugging container:

```
root@gke-nando-1-default-pool-e6acb66d-6lvz:/# apt-get update && apt-get install sysstat
```

# Troubleshooting

In some particular circumstances, it may be necessary to scale the cluster down to 1 node (*unsafe configurations*), remove the data from the scaled-down nodes, and then scale the cluster back again - for example, to force the cluster to “rebootstrap” from node 0:

```
$ kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS
datadir-cluster1-pxc-0	Bound	pvc-76d2deb4-965a-4a4c-a23c-b12f869a00b8	10Gi	RWO	standard
25h					
datadir-cluster1-pxc-1	Bound	pvc-83cce7ad-4691-4d60-87eb-fe853a53068e	10Gi	RWO	standard

```
data: spec:
  allowUnsafeConfigurations: true
  pxc:
    size: 1
```

```
$ kubectl apply -f deploy/cr.yaml
```

```
$ kubectl delete pvc datadir-cluster1-pxc-1 \
datadir-cluster1-pxc-2
```

```
spec:
  allowUnsafeConfigurations:
false
  pxc:
    size: 3
```

```
$ kubectl apply -f deploy/cr.yaml
```

# Troubleshooting

Avoid the restart-on-fail loop by preventing *mysqld* to start on the node:

```
$ kubectl exec -it cluster1-pxc-0 -c logs -- sh -c 'touch /var/lib/mysql/sleep-forever'
```

- Both the *pxc* and *logs* containers have access to the storage with the *datadir*
- Remove the *sleep-forever* file to allow *mysqld* to start again

There are special debug images for the officially supported PXC releases; they are identified by the *-debug* suffix. It is possible to restart the pods with them:

```
spec:
  pxc:
    image: percona/percona-xtradb-cluster:8.0.29-21.1-debug
```

```
$ kubectl apply -f deploy/cr.yaml
```

Extras



# Extras

---

- ~~Deploying with a Helm chart~~
- ~~Installing PXC in multi-namespace mode~~
- ~~Multi-cluster and multi-region deployment~~
- Monitoring with PMM
- Upgrading

# Administrative tasks: monitoring

---

- PMM → <https://docs.percona.com/percona-operator-for-mysql/pxc/monitoring.html>
- Add your PMM server's API key (value) to `deploy/secrets.yaml` as *pmmserverkey*

```
pmm:
  enabled: true
  image: percona/pmm-client:2.32.0
  serverHost: <pmm-server>
#   serverUser: admin
#   pxcParams: "--disable-tablestats-limit=2000"
#   proxysqlParams: "--custom-labels=CUSTOM-LABELS"
  resources:
    requests:
      memory: 150M
      cpu: 300m
```

# Administrative tasks: upgrading the operator

<https://docs.percona.com/percona-operator-for-mysql/pxc/update.html>

- Download the new Custom Resource Definition file (crd.yaml) and the Role-based access control file (rbac.yaml) and apply them

```
$ kubectl apply -f  
https://raw.githubusercontent.com/percona/percona-xtradb-cluster-operator/v1.12.0/deploy/crd.yaml  
  
$ kubectl apply -f  
https://raw.githubusercontent.com/percona/percona-xtradb-cluster-operator/v1.12.0/deploy/rbac.yaml
```

- Apply a patch to the deployment, specifying the image to use

```
$ kubectl patch deployment percona-xtradb-cluster-operator \  
-p '{"spec":{"template":{"spec":{"containers":[{"name":"percona-xtradb-cluster-operator",  
"image":"percona/percona-xtradb-cluster-operator:1.12.0"}]}}}}'
```

- Track the rollout progress with:

```
$ kubectl rollout status deployments percona-xtradb-cluster-operator
```

# Administrative tasks: upgrading PXC

---

- `updateStrategy`: `SmartUpdate`
- `upgradeOptions.apply`:
  - Recommended → 8.0 for new clusters, otherwise preserve the existing branch
  - 8.0-recommended or 5.7-recommended → a way to start with 5.7
  - Latest → automatic upgrades pick the most recent version
  - 8.0-latest or 5.7-latest → a way to be using the latest in 5.7
  - `<version-number>` → specify a specific version of PXC
  - Never or Disabled → disables automatic upgrades

The *SmartUpdate* method relies on `versionServiceEndpoint`  
→ <https://check.percona.com>

\* *manual* and *semi-automatic* update methods reserved for Operator version 1.5.0 and earlier



# Thank you!

[fernando.laudares@percona.com](mailto:fernando.laudares@percona.com)  
[chetan.shivashankar@percona.com](mailto:chetan.shivashankar@percona.com)